

Assignment 2

In-memory KV-Store

By Shabir Abdul Samadh
Oct 16th/17th 2018

Topics

- Key-Value store?
- Hash Functions
- General discussion/description of the assignment
 - `memset()`
 - `memcpy()`
 - Synchronization
 - Use `gdb` / `hexedit`


Key-Value Store

- A data storage paradigm designed for storing, retrieving, and managing associative arrays,
- Different from relational databases where the data has a specific structure
 - With KV-Store, values could be anything
- Data is treated as a single unit
- In memory KV-Store? (eg: Redis)
 - Kept in the system-memory (RAM)
 - Persisted until system reboot

Key	Value
K1	AAA,BBB,CCC
K2	AAA,BBB
K3	AAA,DDD
K4	AAA,2,01/01/2015
K5	3,ZZZ,5623

Hash Functions

- Any **function** that can be used to map data of arbitrary size to data of a fixed size
- The values returned by a hash function are called **hash values**, **hash codes**, **digests**, or simply **hashes**
- Always gives the same hash value for a specific data segment
- One Way functions



INPUT	HASH
This is a test	C7BE1ED902FB8DD4D48997C6452F5D7E509FBCDBE2808B16BCF4EDCE4C07D14E
this is a test	2E99758548972A8E8822AD47FA1017FF72F06F3FF6A016851F45C398732BC50C


Hash Functions

```
unsigned long hash(unsigned char *str) {  
    unsigned long hash = 5381;  
    int c;  
  
    while (c = *str++)  
        hash = ((hash << 5) + hash) + c;  
  
    return (hash > 0) ? hash : -(hash);  
}
```

```
unsigned long hash(unsigned char *str) {  
    unsigned long hash = 0;  
    int c;  
  
    while (c = *str++)  
        hash = c + (hash << 6) + (hash << 16) - hash;  
  
    return (hash > 0) ? hash : -(hash);  
}
```

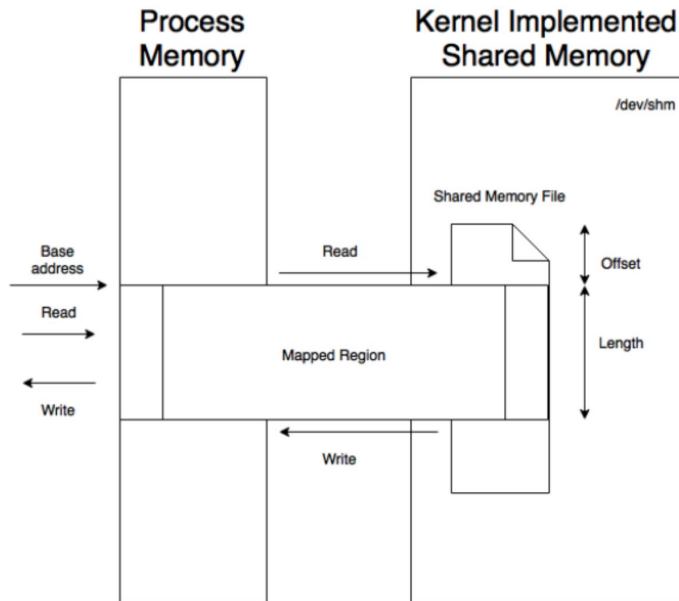
This will give a long number

You have to reduce it to the range of your KVStore



Shared Memory

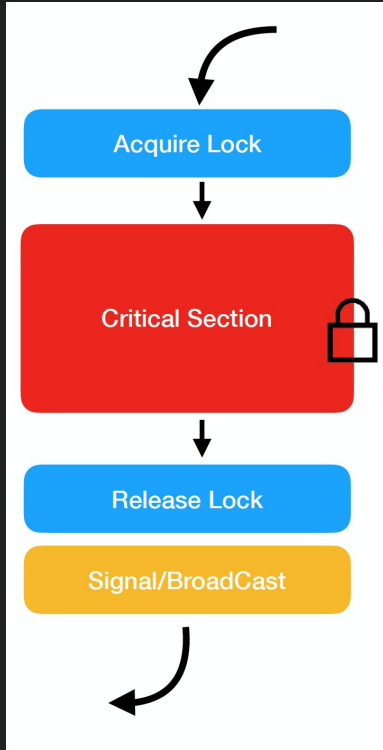
Open SharedMem File
↓
Map it kernel Shared Mem
↓
Read or write
↓
Close the file



`shm_open()`
`ftruncate()`
`mmap()`
`memcpy()`
`munmap()`

read [man page](#) for proper syntax
created under [/dev/shm](#)

Synchronization



`sem_open()`

- **for creation** : requires mode and value
- **just open**

`sem_wait()`

`sem_post()`

`sem_close()`

`sem_unlink()`

Read [man page](#) for proper syntax

Assignment Description

```
int kv_store_create(char *kv_store_name);
```

- open shared memory segment
- `ftruncate()` to set the KVStore size
- `mmap()` to process virtual memory
- initialize bookkeeping information

```
int kv_store_write(char *key, char *value);
```

- open shared memory segment
- `mmap()` to process virtual memory
- calculate hash of key (if using hash)
- store the K-V pair in the right location
- **many values are possible a for single key (store all of them)**

Assignment Description

```
char *kv_store_read(char *key);
```

- open shared memory segment
- `mmap()` to process virtual memory
- calculate hash of the key to get the location
- get to location -----> read value

```
char **kv_store_read_all(char *key);
```

- same as above but remember each key can have multiple values
- read all of them and return

Some key points (there could be more....)

```
int kv_store_create(char *kv_store_name);
```

- check for existence of a KV-Store with same name? **O_EXCL**
- ensure all open fd(s) are close
- ensure to **munmap()**

```
int kv_store_write(char *key, char *value);
```

- ensure exclusive access via **LOCK**
- ensure key & value are within range (32 bytes and 256 bytes)
- if same key with a new value? Write key twice?
- if two keys map to the same location, what to do?
- if there is not enough space, then which value to replace
 - How will you keep track of this value?
 - If that is also not enough?
- clean up everything before leaving the function

Some key points (there could be more....)

```
char *kv_store_read(char *key);
```

- ensure exclusive **LOCK** over writers, but readers can join in
- check if the key exists in the KV-Store
 - if key not found before returning release **LOCKS (cleanup)**
- if there are many records which one to return?

```
char **kv_store_read_all(char *key);
```

- can you call the above method inside this?
 - what if there is a **write()** in between two calls?

Some key points (there could be more....)

General

- exclusive **LOCK** as needed (based on reader or writer)
- always remember to cleanup once done/before returning
 - `free()`
 - `munmap()`
 - `sem_close()` or `sem_post()`
 - `close(fd)`
- Use `calloc()` instead of `malloc()`
 - or `memset()` to `'\0'` **NULL** upon `malloc()`
 - use `memcpy()` to transfer exact amounts of bytes
 - truncate strings with `'\0'` **NULL**
- be cautious of the type of the pointer to memory address

Some key points (there could be more....)

- Tester requires you to implement a method to clear stuff
 - `void kv_delete_db();`
 - `sem_unlink();` -----> all named semaphores
 - `kill_shared_mem();` -----> already provided
 - name your main source file `a2_lib.c`
 - might have to add other source files to “make” if you have more
- If the program is stuck between subsequent calls - maybe semaphore was waited for and not `post()`'ed again.
 - remove it from `/dev/shm`
 - use very unique names for semaphores and KV_Store
 - `comp310_james`, `a2_database` (**not unique!!!**)
- `cat /dev/shm/KV_STORE > some_file`
 - then view it on HexEdit to see what's on it

Synchronization

Writer

```
obtain lock(WRITE_LOCK)
write stuff
release lock(WRITE_LOCK)
```

```
semaphore - WRITE_LOCK
semaphore - READ_LOCK
int variable - READ_COUNT
```

named semaphores are system
level shared

how to share read count?

Reader

```
obtain lock(READ_LOCK)
increment(READ_COUNT)
If (READ_COUNT == 1)
    obtain lock(WRITE_LOCK)
release lock(READ_LOCK)
```

..... readers can come in but not writers

read stuff

..... readers can come in but not writers

```
obtain lock(READ_LOCK)
decrement(READ_COUNT)
If (READ_COUNT == 0)
    release lock(WRITE_LOCK)
release lock(READ_LOCK)
```

Thank You!!!

Have *patience*...

Use **gdb** to debug to check variables & memory segments...

Comment out portions of the tester and try one by one...

Open thread on **myCourses**, *email me* and I will try my best to help

Good Luck!